

Getting good performance from your application

Tuning techniques for serial programs on
cache-based computer systems

Overview

- ❑ Introduction
- ❑ Memory Hierarchy
- ❑ General Optimization Techniques
- ❑
- ❑ Compilers
- ❑ Analysis Tools
- ❑ Tuning Guide

Introduction

Introduction

Moore's Law

- ❑ Popular version:

- ❑ *“CPU speed usually doubles every 18 months.”*

- ❑ More correct version:



- ❑ *“The number of transistors per integrated circuit will double every 18 months.”*

Introduction


Gordon Moore – co-founder of Intel

"I never said 18 months. I said one year, and then two years ... Moore's Law has been the name given to everything that changes exponentially. ... If Gore invented the Internet, I invented the exponential."

- Gordon Moore in an interview (2000)




Introduction




Introduction

Maximum Intel CPU Speed (IA-32) vs Time

Connelly Barnes
Public domain, 2005-11-13



Introduction




Introduction

- ❑ CPU speed usually doubles every 18-24 months (not true any longer – we got other improvements, instead!).
- ❑ Development on the memory side is much slower (~ 4-6 years!).
- ❑ Memory speeds catch up – but also have to serve more cores!
- ❑ New memory technologies do help (e.g. HBM)
- ❑ Something you should have in mind when designing your program!


Motivation for Application Tuning

time flow in a computational task (simplified picture):



→ *time*

new hardware:



ideal: 2x faster computer



reality: 2x faster CPU


code tuning (old hardware):



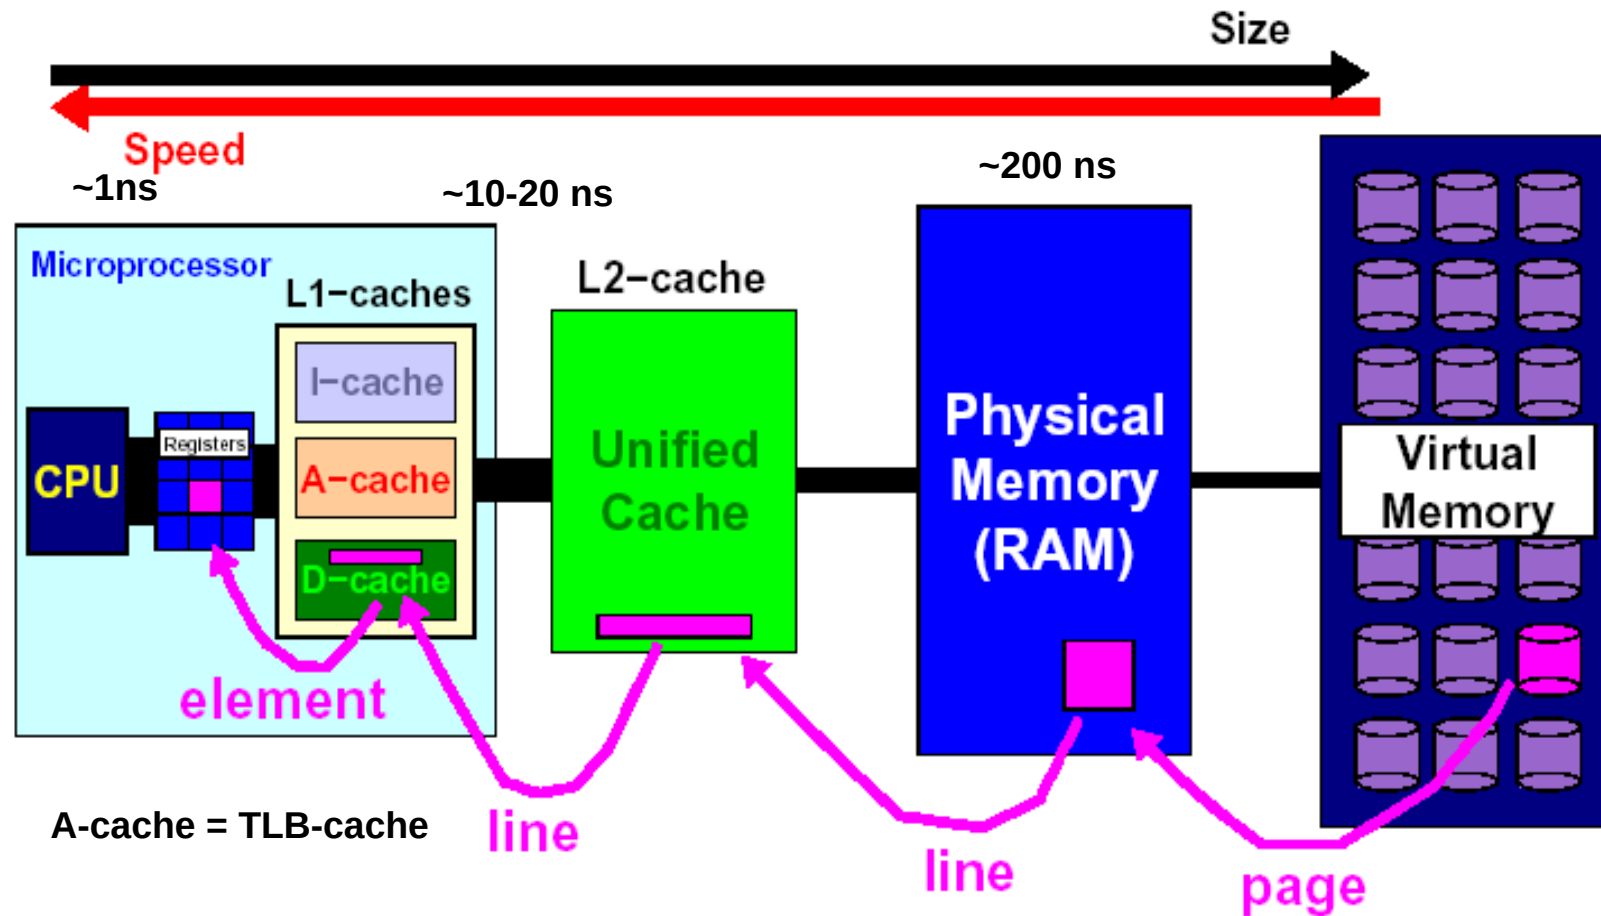
The Memory Hierarchy

The Memory Hierarchy

Intuitive Performance Graph:




The Memory Hierarchy




Memory Optimization:
Keep frequently used data close to the processor

The Memory Hierarchy

Performance is not uniform:




The Memory Hierarchy



Intel Xeon X5550, 2.67 GHz, 8 MB L3 cache (release: Q1/2009)

The Memory Hierarchy



Intel Xeon E5-2660v3, 2.66 GHz, 25 MB L3 (release: Q3/2014)

The Memory Hierarchy

- ❑ Memory plays a crucial role in performance
- ❑ Not accessing memory in the right way will degrade performance on **all** computer systems
- ❑ The extent of degradation will depend on the system
- ❑ Knowledge about the relevant memory characteristics helps to write code that minimizes those problems

Caches – and all that ...


How do those caches work?

Caches

- ❑ Cache memory or cache for short (from French: cacher – to hide): fast buffers that help to hide the memory latency
- ❑ One distinguishes between
 - ❑ data cache
 - ❑ instruction cache
 - ❑ address cache (also called TLB – Translation Lookaside Buffer) – mapping between virtual and physical addresses

Cache Lines

- ❑ To get good performance, optimal use of the caches is crucial
- ❑ The unit of transfer is a “*cache line*”:
 - ❑ linear structure of fixed length (bytes)
 - ❑ fixed starting address in memory



Cache Organisation

Direct Mapped:

- ❑ Each memory address maps onto exactly one line in cache
- ❑ simple and efficient
- ❑ built-in replacement policy
- ❑ easy to scale to larger sizes
- ❑ downside: no control by usage – danger of replacing data that will be needed again soon

Cache Organisation

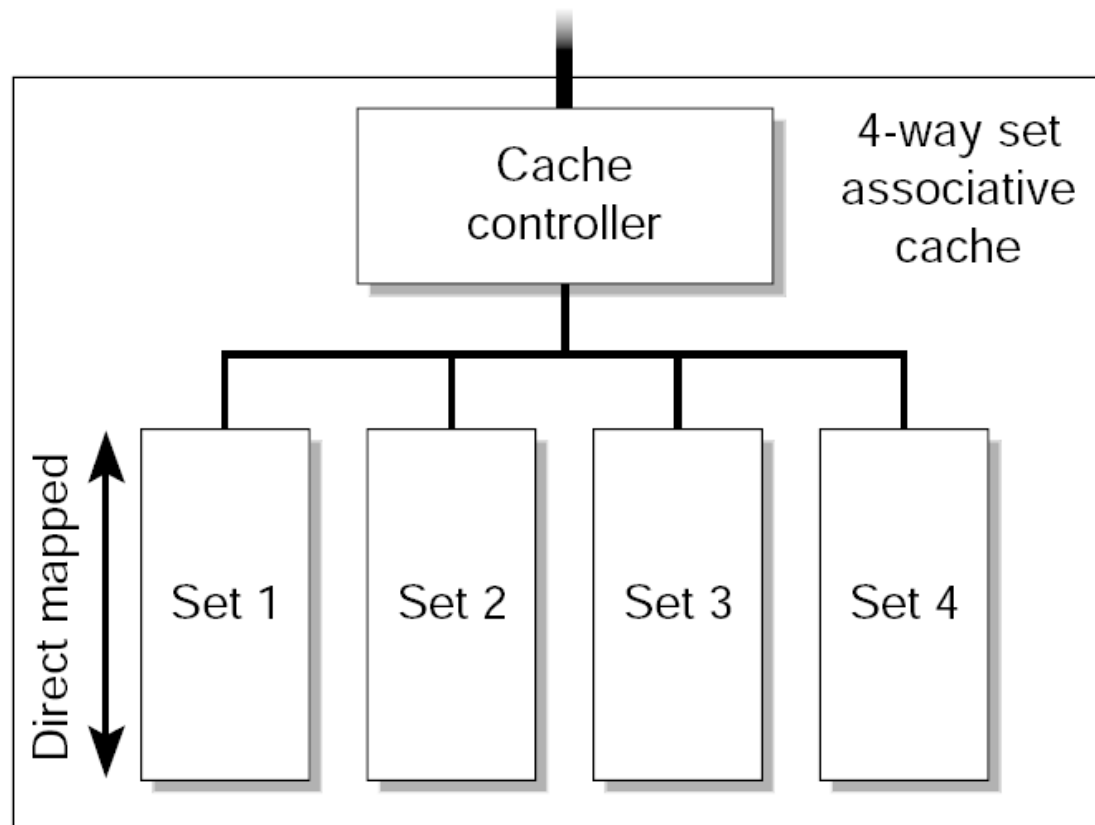
Fully Associative:

- ❑ Every memory address can be mapped anywhere in cache
- ❑ Need to track usage of cache lines
- ❑ Requires a replacement policy, e.g.
 - ❑ *least recent used* (LRU),
 - ❑ *least frequent used* (LFU),
 - ❑ random, etc
- ❑ Doesn't scale well to large sizes
- ❑ Costly design

Cache Organisation


N-way Set Associative:

- ❑ Sets of direct mapped caches:



Memory access

- ❑ Memory has a 1-dimensional linear structure
- ❑ Accessing vector elements:




size of a
cache line


Legend:



vector element



cache miss




cache hit







memory access

Memory access

Accessing 2d arrays in C – row wise:

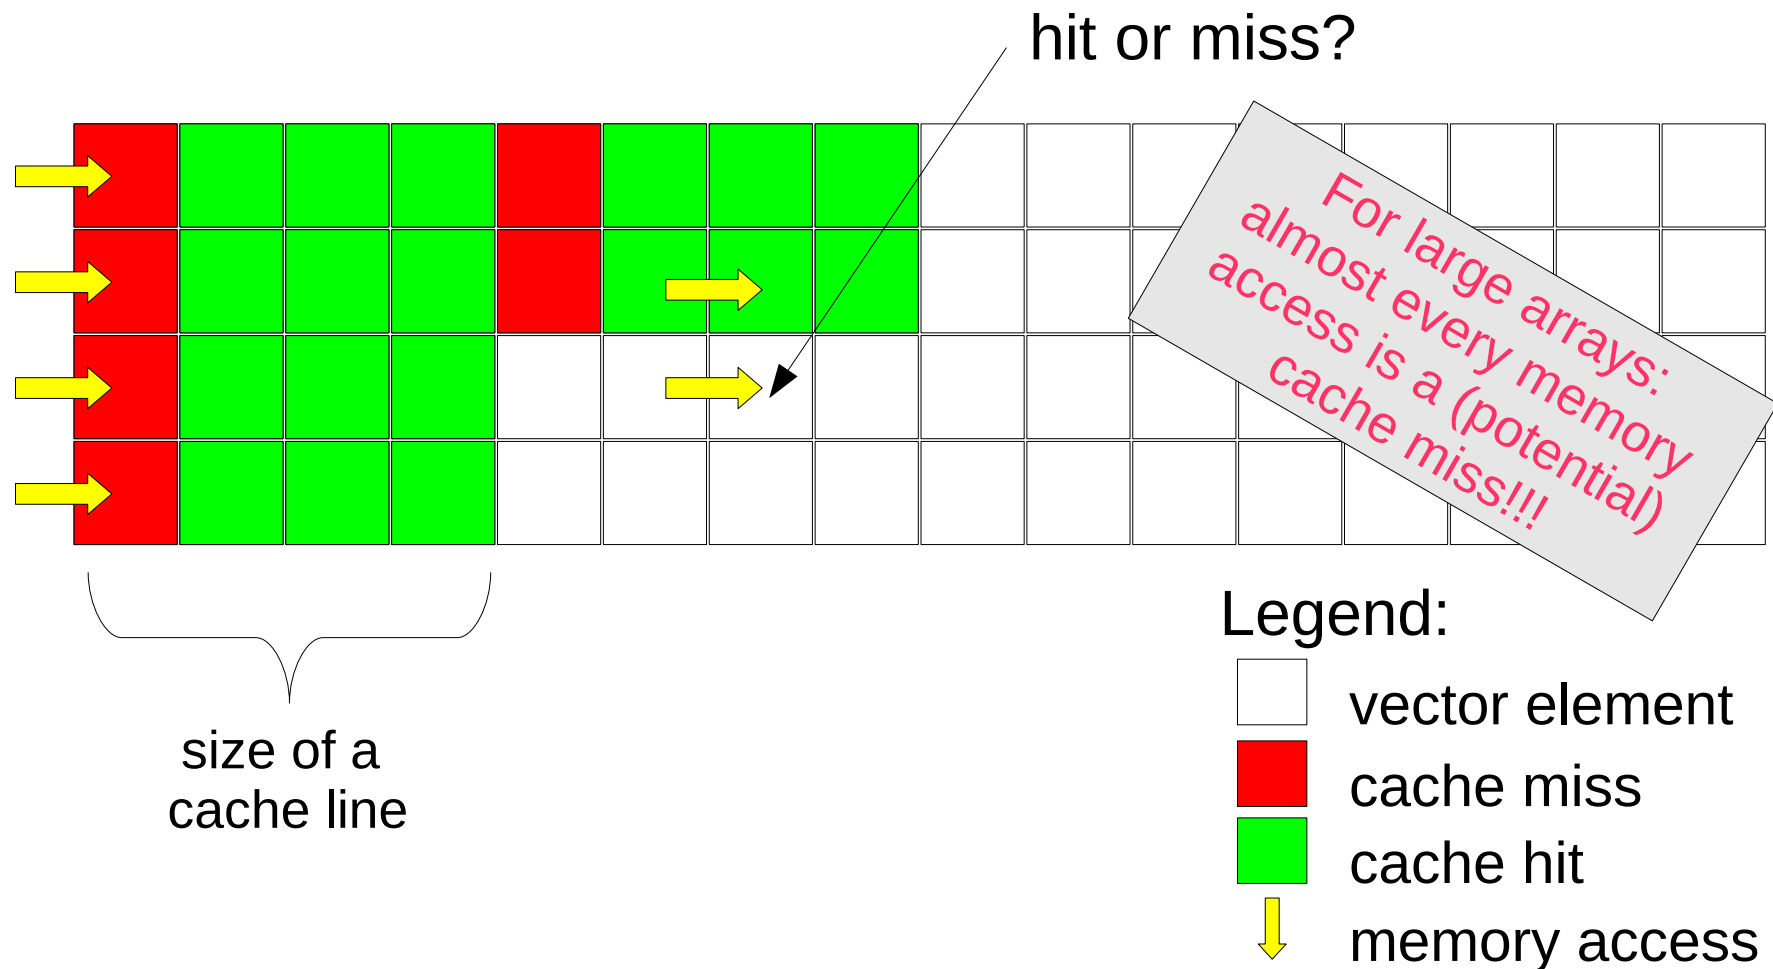


Legend:

-  vector element
-  cache miss
-  cache hit
-  memory access


Memory access

Accessing 2d arrays in C – column wise:



Memory access

Access to multi-dimensional arrays depends on how data is stored:




Bad memory access has a huge impact on performance!!!

The TLB cache

- ❑ the Translation Lookaside Buffer (TLB) translates virtual memory addresses (in your application) to physical addresses
- ❑ also called 'address cache'
- ❑ unit: page – typical size 4kB
- ❑ creation of lookup table is an expensive operation
- ❑ cost: 10 – 100 clock cycles/miss
- ❑ modern CPUs are having advanced TLBs
 - ❑ support for variable page sizes


Memory access – TLB misses

storage order →
good access order



storage order →

bad access order ↓



- = TLB miss
- = D-cache miss
- = Cached elements
- ■ = Virtual memory page

- If the entire matrix fits in the cache, the access pattern *hardly* matters.
- For large (out-of-cache) matrices, the access pattern **does** matter – both data cache and TLB misses

About cache misses

Some simple rules:

- ❑ You cannot avoid cache misses – they are part of the nature of cache-based systems ...
- ❑ ... but you should try to minimize them to get good performance

Cache Line Utilization

Two key rules: **Maximize ...**

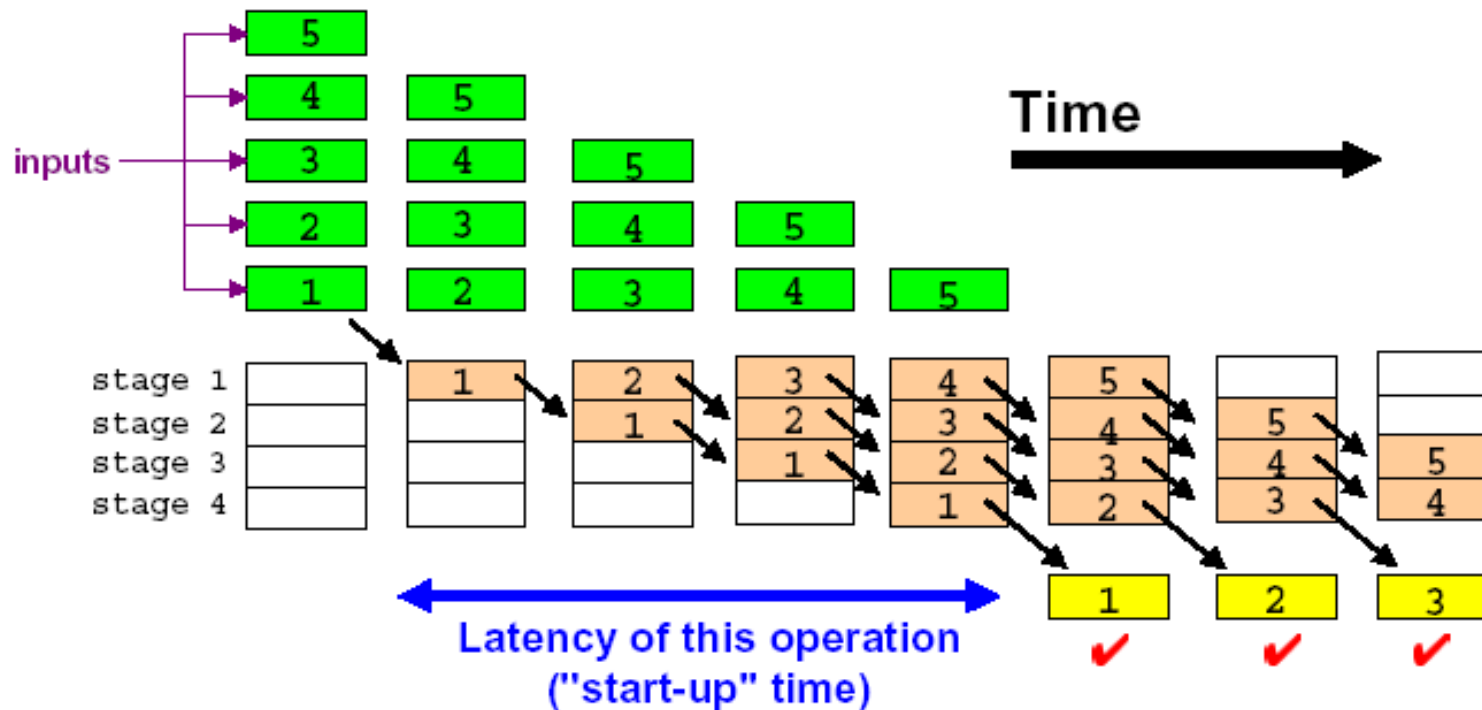
- ❑ **Spatial locality** \Rightarrow Use all data in one cache line
 - ❑ depends on storage layout
 - ❑ depends on access patterns
 - ❑ stride = 1 is good
 - ❑ random access is really bad
- ❑ **Temporal locality** \Rightarrow Re-use data in a cache line
 - ❑ depends on algorithm used

Some Terminology

Terminology: Pipelining

stage 1
stage 2
stage 3
stage 4

Let's assume we have an operation that takes 4 stages per iteration



Rule of thumb: keep the pipeline filled for good performance!

Terminology: Superscalar (or ILP)

❑ *N-way superscalar:*

- *Execute N instructions at the same time*

❑ *This is also called Instruction Level Parallelism (ILP)*

	slot 1	slot 2	slot 3	slot 4	
cycle 1					4-way superscalar
cycle 2	not used				3-way superscalar
cycle 3		not used	not used		2-way superscalar
cycle 4	not used			not used	2-way superscalar
cycle 5			not used		3-way superscalar

- ❑ *The hardware has to support this, but it is up to the software to take advantage of it*
- ❑ *Often there are restrictions which instructions can be "bundled"*
- ❑ *These are documented in the Architecture Reference Manual for the microprocessor*

Terminology: FMA

One variant of ILP for floating point operations:

- ❑ FMA – Fused Multiply Add
- ❑ special instruction, that calculates
$$a + (b * c)$$
in one operation, with a single rounding
- ❑ faster
- ❑ ... but the result might slightly differ, due to the single rounding!


Terminology: Vectorization (SIMD)

- ❑ Vectorization is the process of transforming a scalar operation that acts on a single data element at a time (**Single Instruction Single Data – SISD**) to an operation that acts on multiple data elements at a time (**Single Instruction Multiple Data – SIMD**).
- ❑ **SIMD** units are hardware arithmetic vector units that can perform the same operation on multiple data points simultaneously by using vector registers.

Terminology: Vectorization (SIMD)

□ Vector types


- 128 bit: SSE = Streaming SIMD Extension (1999)



- 256 bit: AVX = Advanced Vector Extension (2011)



- 512 bit: AVX-512 (2013)



Latency and Bandwidth

Latency:

- ❑ the time it takes from the initiation of an action till you have the first result
- ❑ unit: time

Bandwidth:

- ❑ how many
 - ❑ actions can be carried out,
 - ❑ results can be obtained,within a given time
- ❑ unit: #/time

General Optimization Techniques

Optimization Techniques - Overview

- ❑ Most optimization techniques are “loop based”
- ❑ Loop based optimizations:
 - ❑ Interchange
 - ❑ Fission and Fusion
 - ❑ Unrolling
 - ❑ Blocking

Optimization Techniques - Overview

- ❑ Designing your data structures the “right way” can also be important
- ❑ Other techniques:
 - ❑ De-vectorization
 - ❑ Stripmining



Loop based optimizations

Coding style: array indexing

- ❑ To apply safe transformations, the compilers have to analyze data dependencies in a loop
- ❑ Explicit expressions will help the compilers to do a good job in loop optimization

Good

```
for(i=0; i<m; i++)  
  for (j=0; j<n; j++)  
    .. a[i][j] ..
```

() Reasonable*

```
for(i=0; i<m; i++)  
  for (j=0; j<n; j++)  
    .. a[i*n+j] ..
```

Bad

```
k = 0;  
for(i=0; i<m; i++)  
  for (j=0; j<n; j++)  
    .. a[k++] ..
```

Caution

```
for(i=0; i<m; i++)  
  for (j=0; j<n; j++)  
    .. a[indx[i][j]] ..
```

(*) harder to read for a human, but might be better for the compiler!

Loop Interchange

```
DO I = 1, M
  DO J = 1, N
    A(I, J) = B(I, J) + C(I, J)
  END DO
END DO
```

Interchange
loops

```
DO J = 1, N
  DO I = 1, M
    A(I, J) = B(I, J) + C(I, J)
  END DO
END DO
```

- ❑ The matrices are accessed over the second dimension first
- ❑ This is the wrong order in Fortran
- ❑ A loop interchange solves the problem
- ❑ In C, the situation is reversed:
 - ❑ row access is okay
 - ❑ column access is bad

Loop Fission

```
for (j=0; j<n; j++)  
{  
    c[j] = exp(j/n);  
    for (i=0; i<m; i++)  
        a[i][j]=b[i][j]+d[i]*e[j];  
}
```

- ♦ Access on arrays 'a' and 'b' is bad
- ♦ We can not simply interchange the loops
- ♦ Fission/splitting is the solution

Fission

This loop can now also be vectorized

Interchange loops for better performance

```
for (j=0; j<n; j++)  
    c[j] = exp(j/n);
```

New loop created

```
for (j=0; j<n; j++)  
    for (i=0; i<m; i++)  
        a[i][j]=b[i][j]+d[i]*e[j];
```

Loop Fusion

```
for (i=0; i<n; i++)  
    a[i] = 2 * b[i];
```

```
for (i=0; i<n; i++)  
    c[i] = a[i] + d[i];
```

Fusion




Note that it is possible to apply fusion to loops with (slightly) different boundaries

In such a case, some iterations will have to be 'peeled' off

- ♦ Assume that 'n' is large
- ♦ In the second loop, a[i] will no longer be in the cache
- ♦ Fusing the loops will ensure a[i] is still in the cache when needed

```
for (i=0; i<n; i++)  
{  
    a[i] = 2 * b[i];  
    c[i] = a[i] + d[i];  
}
```

Fission and Fusion – Summary



Fission

- ✓ Reduce register pressure
- ✓ Enable loop interchange
- ✓ Isolate dependencies
- ✓ Increase opportunities for optimization (e.g. vectorization of intrinsics)

Fusion

- ✓ Reduce cache reloads
- ✓ Increase Instruction Level Parallelism (ILP)
- ✓ Reduce loop overhead

(Inner) Loop Unrolling

Through unrolling, the loop overhead ('book keeping') is reduced

```
for (i=0; i<n; i++)
    a[i] = b[i] + c[i];
```

*Loop is unrolled
with a factor of 4*

```
for (i=0; i<n; i+=4)
{
    a[i  ] = b[i  ] + c[i  ];
    a[i+1] = b[i+1] + c[i+1];
    a[i+2] = b[i+2] + c[i+2];
    a[i+3] = b[i+3] + c[i+3];
}
<clean-up loop>
```

```
Loads      : 2
Stores     : 1
FP Adds    : 1
I=I+1
Test I < N ?
Branch
Addr. incr: 3
```


Work: 4
Overhead: 6

```
Loads      : 8
Stores     : 4
FP Adds    : 4
I=I+4
Test I < N ?
Branch
Addr. incr: 3
```

Work: 16
Overhead: 6

Note: the amount of addressing needed in reality is less

Outer Loop Unrolling



```
for (i=0; i<m; i++)
  for(j=0; j<n; j++)
  {
    a[i] += b[i][j] * c[j];
  }
```

```
for (i=0; i<m; i+=4)
  for(j=0; j<n; j++)
  {
    a[i  ] += b[i  ][j] * c[j];
    a[i+1] += b[i+1][j] * c[j];
    a[i+2] += b[i+2][j] * c[j];
    a[i+3] += b[i+3][j] * c[j];
  }
<clean-up loop>
```

♦ Advantage:

- *c[j]* is re-used 3 more times (temporal locality)
- ♦ Deeper unrolling, say 8, requires more fp registers (17 instead of 9), but improves re-use of *c[j]*

Outer Loop Unrolling – how to

```
for (i=0; i<m; i++)
  for(j=0; j<n; j++)
    a[i] += b[i][j] * c[j];
```

*Outer loop
unrolling*

Unroll and Jam

```
for (i=0; i<m-m%4; i+=4)
  for(j=0; j<n; j++)
  {
    a[i  ] += b[i  ][j] * c[j];
    a[i+1] += b[i+1][j] * c[j];
    a[i+2] += b[i+2][j] * c[j];
    a[i+3] += b[i+3][j] * c[j];
  }
for (i=m-m%4; i<m; i++)
  for(j=0; j<n; j++)
    a[i] += b[i][j] * c[j];
```

clean-up loop

```
for (i=0; i<m-m%4; i+=4)
{
  for(j=0; j<n; j++)
    a[i  ] += b[i  ][j] * c[j];
  for(j=0; j<n; j++)
    a[i+1] += b[i+1][j] * c[j];
  for(j=0; j<n; j++)
    a[i+2] += b[i+2][j] * c[j];
  for(j=0; j<n; j++)
    a[i+3] += b[i+3][j] * c[j];
}
for (i=m-m%4; i<m; i++)
  for(j=0; j<n; j++)
    a[i] += b[i][j] * c[j];
```

clean-up loop

*Jam the loops
together again*

Loop unrolling – structure

```
for (i=0; i<n; i++)
{
    ... [i] ...
}
```

```
DO I = 1, N
    ... (I) ...
END DO
```

Loop unroll factor
is "unroll"

```
for(i=0; i<n-n%unroll; i+=unroll)
{
    ... [i] ...
    ... [i+1] ...
    ... [i+2] ...

    ... [i+unroll-1] ...
}
```

Unrolled Loop

```
DO I = 1, N-mod(N,unroll), unroll
    ... (I) ...
    ... (I+1) ...
    ... (I+2) ...
```

```
    ... (I+unroll-1) ...
END DO
```

Cleanup Loop

```
for(i=n-n%unroll; i<n; i++)
{
    ... [i] ...
}
```

```
DO I = N-mod(N,unroll)+1, N
    ... (I) ...
END DO
```

Loop Unrolling – Summary

- ❑ More than one iteration per loop pass
- ❑ Inner loop unrolling:
 - ❑ reduce loop overhead
 - ❑ better instruction scheduling
- ❑ Outer loop unrolling:
 - ❑ improve cache line usage (spatial locality)
 - ❑ re-use data (temporal locality)
- ❑ Disadvantages:
 - ❑ more registers needed, clean-up code required


Loop Unrolling – Compilers

- ❑ compilers do (usually) a good job in loop unrolling – leave it to the compiler
- ❑ some compiler need to be told to do loop unrolling – not part of standard optimization, always check if in doubt
- ❑ there are options to control the unroll depth
 - ❑ gcc: `-funroll-loops --params max-unroll-times=n`
 - ❑ clang: `-funroll-loops --unroll-count=n`
 - ❑ Intel: `-unroll[n]` (0: disable loop unrolling)
- ❑ or compiler specific pragmas (see docs)

Loop Blocking – 1

Transposing a matrix

```
for (j=0; j<n; j++)  
  for (i=0; i<n; i++)  
    b[j][i] = a[i][j];
```




- ♦ *Loop interchange will not help here:*
 - *Role of 'a' and 'b' will only be interchanged*
- ♦ *Change of programming language won't help either*
- ♦ *Unrolling the i-loop can be beneficial, but requires more registers and doesn't address TLB-misses*
- ♦ *Loop blocking achieves good memory performance, without the need for additional registers*

Loop Blocking – 2

Blocking and interchanging the I-loop

```
for(i1=0; i1<n; i1+=nbi)
  for (j=0; j<n; j++)
    for (i2=0; i2<MIN(n-i1,nbi); i2++)
      b[j][i1+i2] = a[i1+i2][j];
```



- ♦ *Parameter 'nbi' is the blocking size*
- ♦ *Should be chosen as large as possible*
- ♦ *Actual value depends on the cache to block for:*
 - ✓ L1-cache
 - ✓ L2-cache
 - ✓ TLB
 - ✓

Fortran

```
do i = 1, n
  do i1 = 1, n, nb1
    do i2 = 0, min(n-i1+1, nb1) - 1
```

Loop Blocking – Summary

- ❑ Powerful technique to improve:
 - ❑ memory access (spatial locality)
 - ❑ data re-use (temporal locality)
- ❑ Preserves portability – but blocking size depends on:
 - ❑ cache type/level/capacity (hardware)
 - ❑ data requirements

Loop Blocking – Summary

Recommendations:

- ❑ choose blocking size as large as possible
- ❑ leave space for other data
- ❑ parameterize cache characteristics, especially size
- ❑ How many loops should I block?
 - ❑ depends on the data access pattern, i.e. look for load and store operations

Tricked by the compiler

Fortran code example: long and bulky loop

```
DO I1=1,NAT(IMT)
  IA1=IM1+(I1-1)*3
  DO I2=1,NAT(JMT)
    IA2=IM2+(I2-1)*3
    ... statements removed ...
    DX(1)=XNOW(IMT,IA1+1)-XNOW(JMT,IA2+1)
    DX(2)=XNOW(IMT,IA1+2)-XNOW(JMT,IA2+2)
    DX(3)=XNOW(IMT,IA1+3)-XNOW(JMT,IA2+3)
    CX(1)=CM1(1)-CM2(1)
    CX(2)=CM1(2)-CM2(2)
    CX(3)=CM1(3)-CM2(3)
    ... statements removed ...
  ENDDO
ENDDO
```

Independent of loop indices!!!

Moved the 3 lines above the DO loops – no improvement!
Compiler had done this already!

Only the programmer knows ...

```

subroutine do_calc(...)
real(8), dimension(N,M,O,P) :

!---- data initialization
r = 0.0d0; s = 0.0d0; t = 0

select case(calc_type)
  case(most_of_the_time)
    ...
    r(i,j,k,l) = r(i,j,k,l)
    s(i,j,k,l) = s(i,j,k,l)
    ...

  case(rare_event)
    ...
    r(i,j,k,l) = r(i,j,k,l)
    s(i,j,k,l) = s(i,j,k,l)
    t(i,j,k,l) = t(i,j,k,l)
    ...
end select

```

```

subroutine do_calc(...)
real(8), dimension(N,M,O,P):: r,s,t

!---- data initialization
r = 0.0d0; s = 0.0d0

select case(calc_type)
  case(most_of_the_time)
    ...
    r(i,j,k,l) = r(i,j,k,l) + ...
    s(i,j,k,l) = s(i,j,k,l) + ...
    ...

  case(rare_event)
    t = 0.0d0
    ...
    r(i,j,k,l) = r(i,j,k,l) + ...
    s(i,j,k,l) = s(i,j,k,l) + ...
    t(i,j,k,l) = t(i,j,k,l) + ...
    ...
end select

```

Data structure design

Access your data in the right way

Data structure design

- ❑ “Good advice” from a HPC tutorial: Use data structures to avoid too many index calculations
- ❑ Example: particle simulation in 3D with x, y, z coordinates and some other information about each particle, e.g. distance to origin, particle type
 - ❑ `x[i], y[i], z[i], dist[i], ptype[i]`
 - ❑ turn this into a data structure ...

Data structure design

❑ Particle data structure

```
typedef struct particle {  
    double x, y, z;  
    char ptype;  
    double dist;  
} particle_t;
```

❑ Is this a good idea?

Data structure design

- ❑ Answer: It depends ...
 - ❑ ... on problem size
 - ❑ ... how you access the data
 - ❑ ... and how often
 - ❑ ... cache, CPU, etc.
- ❑ Example: program with 2 functions/routines
 - ❑ calc() - accesses all parts of particle_t
 - ❑ re-use() - accesses particle.dist only
 - ❑ usage ratio of both functions is 1:1

Data structure design

Advantages of the data structure:

- ❑ `particle_t *p;` (and then allocate N)
- ❑ easy to pass data in function calls
 - ❑ `func(p, N);`
 - ❑ ... VS `func(x, y, z, ptype, dist, N);`
- ❑ flexible: I can add new elements to `particle_t`, without having to change the function interfaces/prototypes

Data structure design

Downside of the data structure:

- ❑ memory access is no longer optimal
 - ❑ neighbouring elements of the same type, e.g. 'x' are no longer 'stride 1':
 - ❑ $x[i] \rightarrow x[i+1]$: stride 1
 - ❑ $p[i].x \rightarrow p[i+1].x$: ~stride 5 (in this example)
 - ❑ no good cache line usage
- ❑ compiler cannot optimize loops

Data structure design

Get the best of both worlds!

- ❑ instead of using an “array of struct” (AOS),
- ❑ we can create a “struct of arrays” (SOA):

```
typedef struct particle {  
    double *x, *y, *z;  
    char    *ptype;  
    double *dist;  
} particle_t;
```

Data structure design

Downside of the SOA:

- ❑ memory allocation is more complicated
 - ❑ one call to `malloc()` for each element
- ❑ need to change code:
 - ❑ `p[i].x -> p.x[i]`
 - ❑ ... and function prototypes: `f(*p, N) -> f(p, N)`
- ❑ but that's a 'one time effort'

Data structure design

Advantages of SOA outweigh this by far:

- ❑ memory access is optimal again
- ❑ neighbouring elements of the same type, e.g. 'x' are again 'stride 1':
 - ❑ `p.x[i] -> p.x[i+1]` : stride 1
 - ❑ better cache line usage
 - ❑ access to single components, e.g. `p.dist`, do not need to load 'all the rest'
- ❑ compiler can optimize loops better

Data structure design

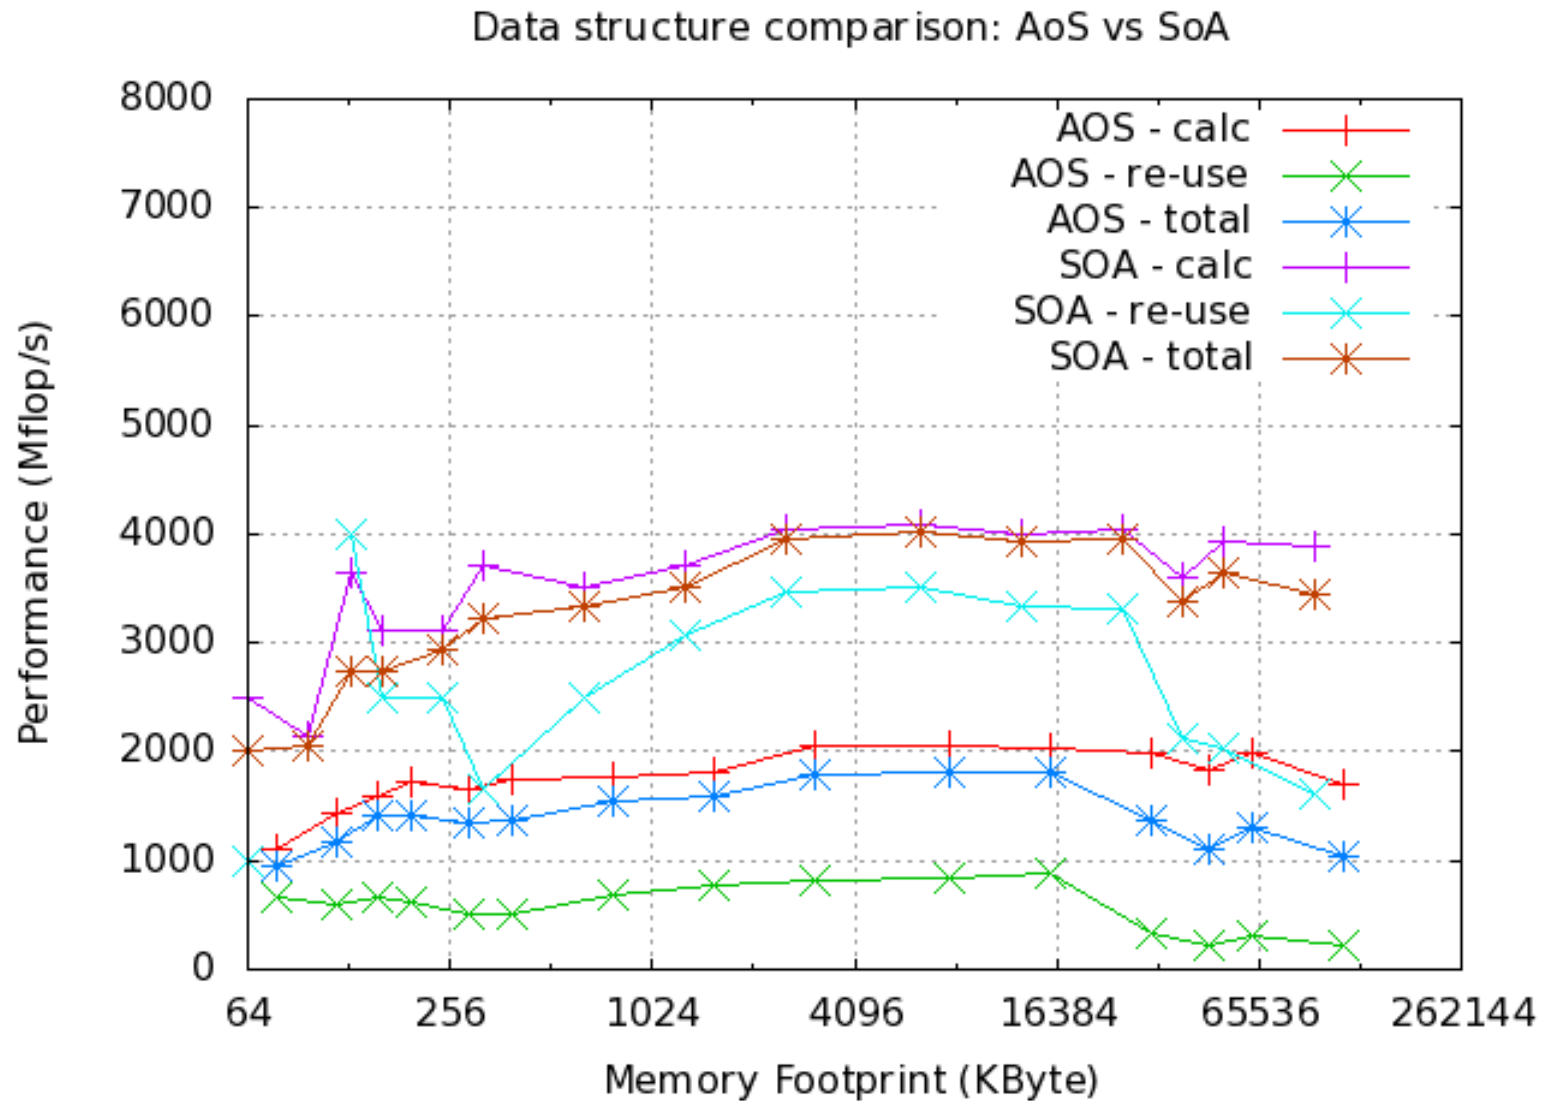
- ❑ Example: program with 2 functions/routines
 - ❑ calc() - accesses all parts of particle_t
 - ❑ re-use() - accesses particle.dist only
 - ❑ usage ratio of both functions is 1:1

Data structure design

Comparison – two versions:


- ❑ *aos* – use ‘array of struct’
- ❑ *soa* – use ‘struct of arrays’
- ❑ different problem sizes:
 - ❑ number of particles

Data structure design




XeonE5-2660v3 25MB L3 cache

Data structure design




XeonE5-2660v3 25MB L3 cache

Data structure design



XeonX5550 8MB L3 cache

Data structure design



XeonX5550 8MB L3 cache

Data structure lab (aka tune_labs)

- ❑ download the ZIP file from DTU Learn
- ❑ read the instructions
 - ❑ do parts I and II now
 - ❑ part III needs more information (after lunch)
- ❑ Goals:
 - ❑ design and do performance experiments
 - ❑ extract the relevant data – and plot it

Performance experiments – how to

Think about the ‘design’ of your experiment!


- ❑ What can I measure directly?
 - ❑ Almost always: time!
 - ❑ Time is not always a “good” quantity
- ❑ Adding ‘a priori’ knowledge: for a given problem size we typically know the ...
 - ❑ number of iterations
 - ❑ number of operations (flops, loads, stores)
- ❑ allows us to calculate Mflop/s, iter/s, bandwidth, ...

Performance experiments – how to


Think about the ‘design’ of your experiment!

- ❑ “Where” should I measure?
 - ❑ your measurement should support a model (or an ‘idea’), e.g. ‘performance vs memory footprint’
 - ❑ make use of prior knowledge, e.g. the memory layout of the hardware (L1,...,L3 cache sizes, and latencies/bandwidths)
 - ❑ choose your measuring points, e.g. memory footprint, to map the “interesting areas” in the “performance landscape”

Performance experiments – how to



Performance experiments – how to



Performance experiments – how to

Best practice:

- ❑ isolate the “kernel” you want to benchmark (e.g. in a function)
- ❑ do a pre-measurement call (“warm up”)
- ❑ repeat the measurements to get a “reasonable” run time (e.g. 2-3 secs), take the average
- ❑ make sure to use the same system (CPU type), to get comparable results
- ❑ reduce the “system noise”, i.e. other processes, users, etc – use the batch system

End of lecture 1