

Getting good performance from your application

Tuning techniques for serial programs on
cache-based computer systems

De-vectorization

- ❑ Cache space and bandwidth are scarce resources
- ❑ Compilers know this – but sometimes they have to store data that does not need to be stored.
- ❑ This impacts:
 - ❑ bandwidth
 - ❑ cache capacity
 - ❑ instruction scheduling

De-vectorization

- ❑ A typical problem with scratch data stored in vectors
- ❑ Difficult/impossible for the compiler to detect
- ❑ Depends on coding style

De-vectorization – Example

```
COMMON /SCRATCH/TMP (N)

DO I = 1, N
    TMP (I) = ...
    :
    ... = TMP (I)
END DO
    :
    :

DO I = 1, N
    TMP (I) = ..
    :
END DO
```

- ❑ *Because $TMP()$ is global, the compiler has to store it in the first loop*
- ❑ *In the second loop, $TMP()$ is overwritten, but the compiler will most likely not see this*
- ❑ *The programmer may know that $TMP()$ is a scratch array only*

De-vectorization – Solutions

Array TMP needed later on:

```

REAL  T1, TMP (N)

DO I = 1, N
    T1 = ...
    :
    ... = T1
END DO
    :
    .

DO I = 1, N
    TMP (I) = ..
    :
END DO
  
```

Array TMP not needed later on:

```

REAL  T1

DO I = 1, N
    T1 = ...
    :
    ... = T1
END DO
    :
    .

DO I = 1, N
    T1 = ..
    :
END DO
  
```

Use of local variables in loops

```
double a[N];  
double b[N];  
double x;  
for(int i=0; i<N; i++) {  
    x = f(a[i]);  
    b[i] = g(x);  
}
```

```
double a[N];  
double b[N];  
  
for(int i=0; i<N; i++) {  
    double x = f(a[i]);  
    b[i] = g(x);  
}
```

This version can “sometimes” give compilers better opportunities to optimize the loop!

Stripmining


- ❑ *Large loops are difficult to optimize*
- ❑ *Especially the **register allocation** in the compiler has a hard time and can get confused*
- ❑ *Splitting the loop into smaller loops may improve performance*
- ❑ *However, this may cause scalars (local to the loop) to be replaced by vectors*
- ❑ *On very large loops this will increase memory usage notably*
- ❑ *Through **stripmining** memory usage can be kept under control*

```
DO I = 1, LONG  
  X(I) = ...  
  A = ...  
  Y(I) = A + ...  
END DO
```

Split loop in
two parts

```
DO I = 1, LONG  
  X(I) = ...  
  VA(I) = ...  
END DO  
  
DO I = 1, LONG  
  Y(I) = VA(I) + ...  
END DO
```

Stripmining – Code structure



Best practice

It is up to you to write code such that the compiler can find opportunities for optimization:

- ❑ Write efficient, but clear code
- ❑ Avoid very “fat” (bulky) loops
- ❑ Design your data structures carefully
- ❑ Minimize global data

Best practice

- ❑ Branches:
 - ❑ simplify where possible
 - ❑ try to split the branch part out of the loop
- ❑ Avoid function calls in loops (use inlining)
 - ❑ the compiler will typically inline code, if it has access to it at compile time (same source file)
 - ❑ it is also possible to do this at the linking stage (will be covered later)
- ❑ Leave the low level details to the compiler

Summary

- ❑ Most tuning techniques presented here are generic, i.e. they (probably/hopefully) improve your code on all cache based systems.
- ❑ The *tuning parameters* may be different, though, since they depend on the underlying hardware:
 - ❑ cache sizes and levels
 - ❑ prefetchand your problem's *memory footprint*
- ❑ Use the *best* compiler available on your platform.

How compilers work

Compilers: overview

❑ Compiler Components

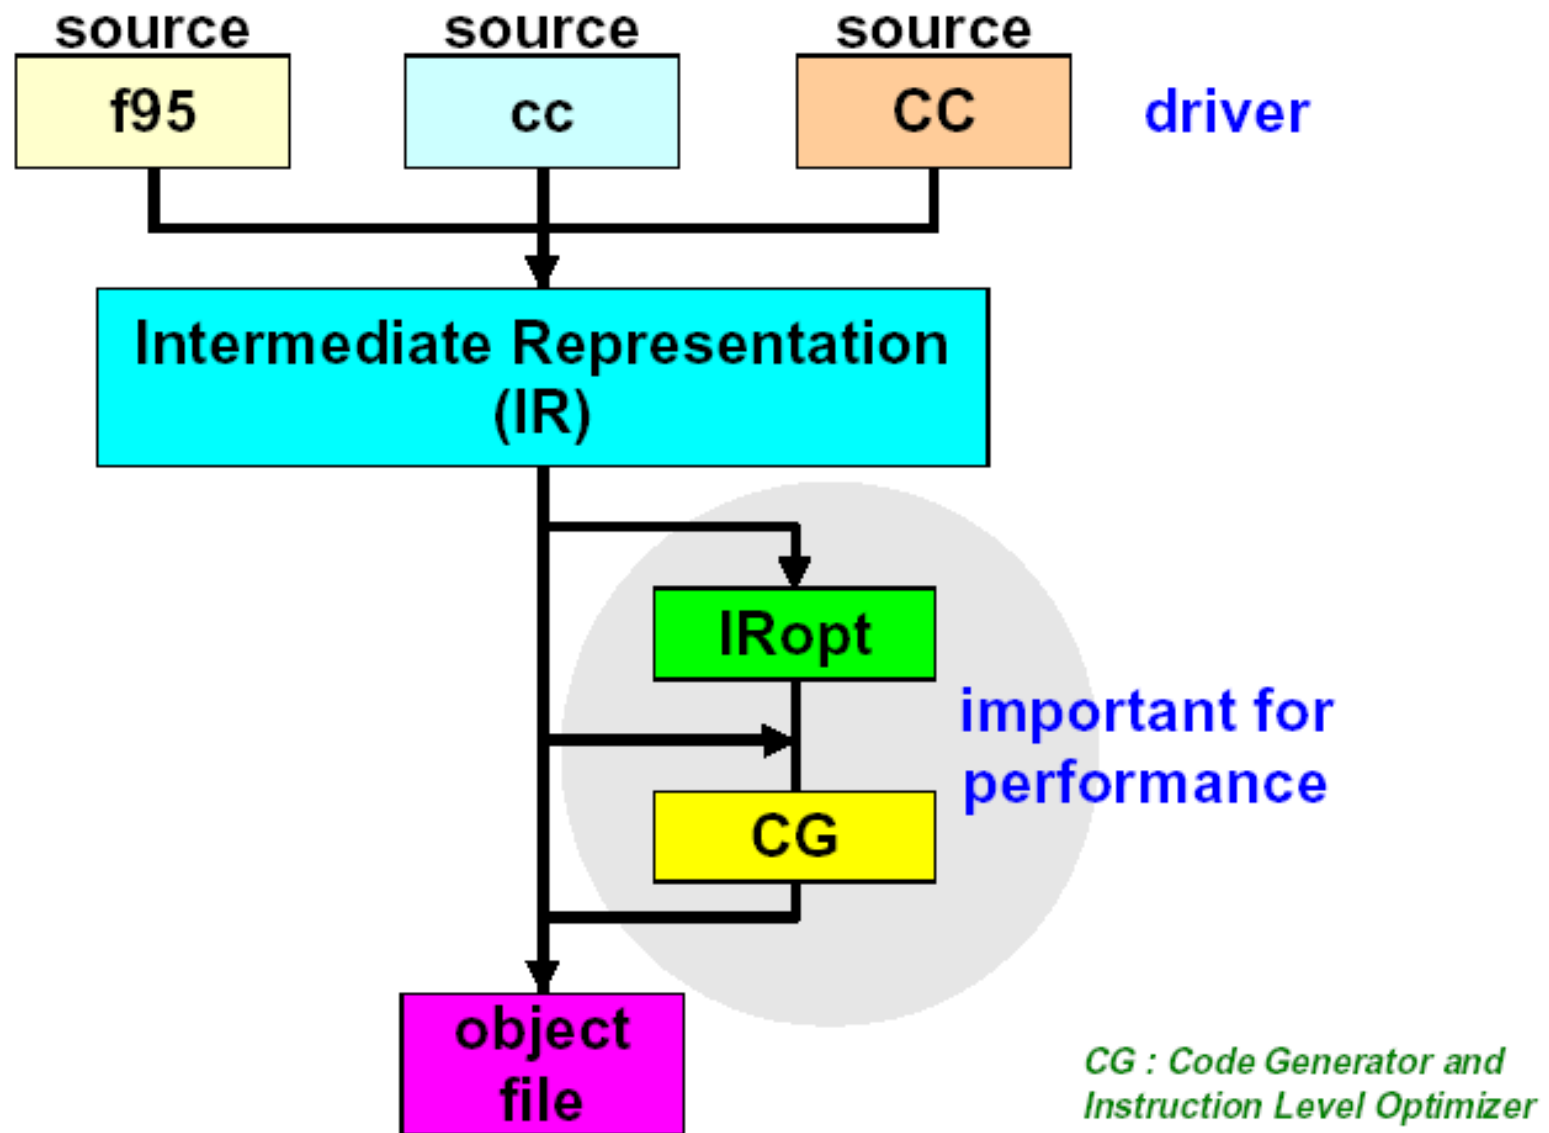
- ❑ compilers are not single programs, but consist of a whole toolchain
- ❑ using Oracle Studio as an example here, for illustration purposes

❑ Compiler Options

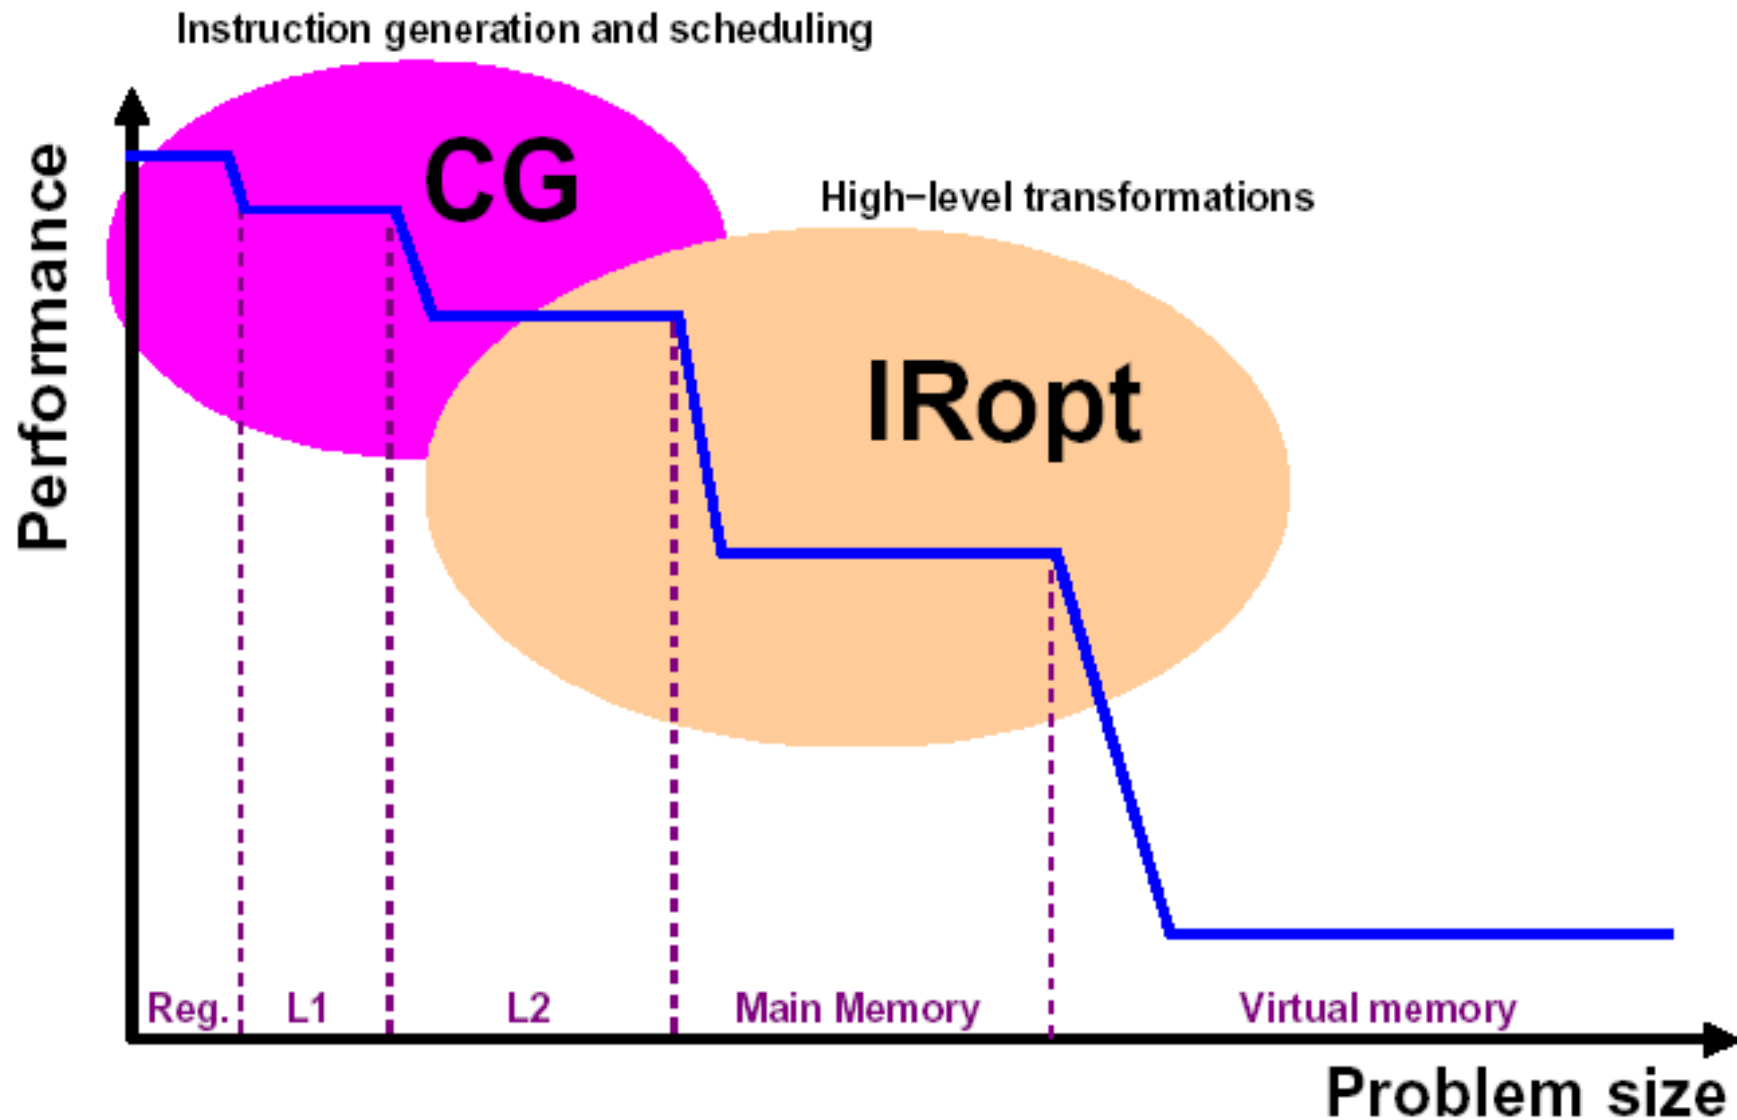
- ❑ minimal set of optimization options ...
- ❑ ... and the more detailed view

❑ Some specific examples

Oracle Studio: Compiler toolchain



Oracle Studio: Who does what?



Oracle Studio: Compiler Options

In general, one gets very good performance by just using 2 options for compiling and linking:

`-g -fast`

For specific x86_64-processors (cross-compiling):

Intel Sandy Bridge: `-g -fast -xchip=sandybridge`

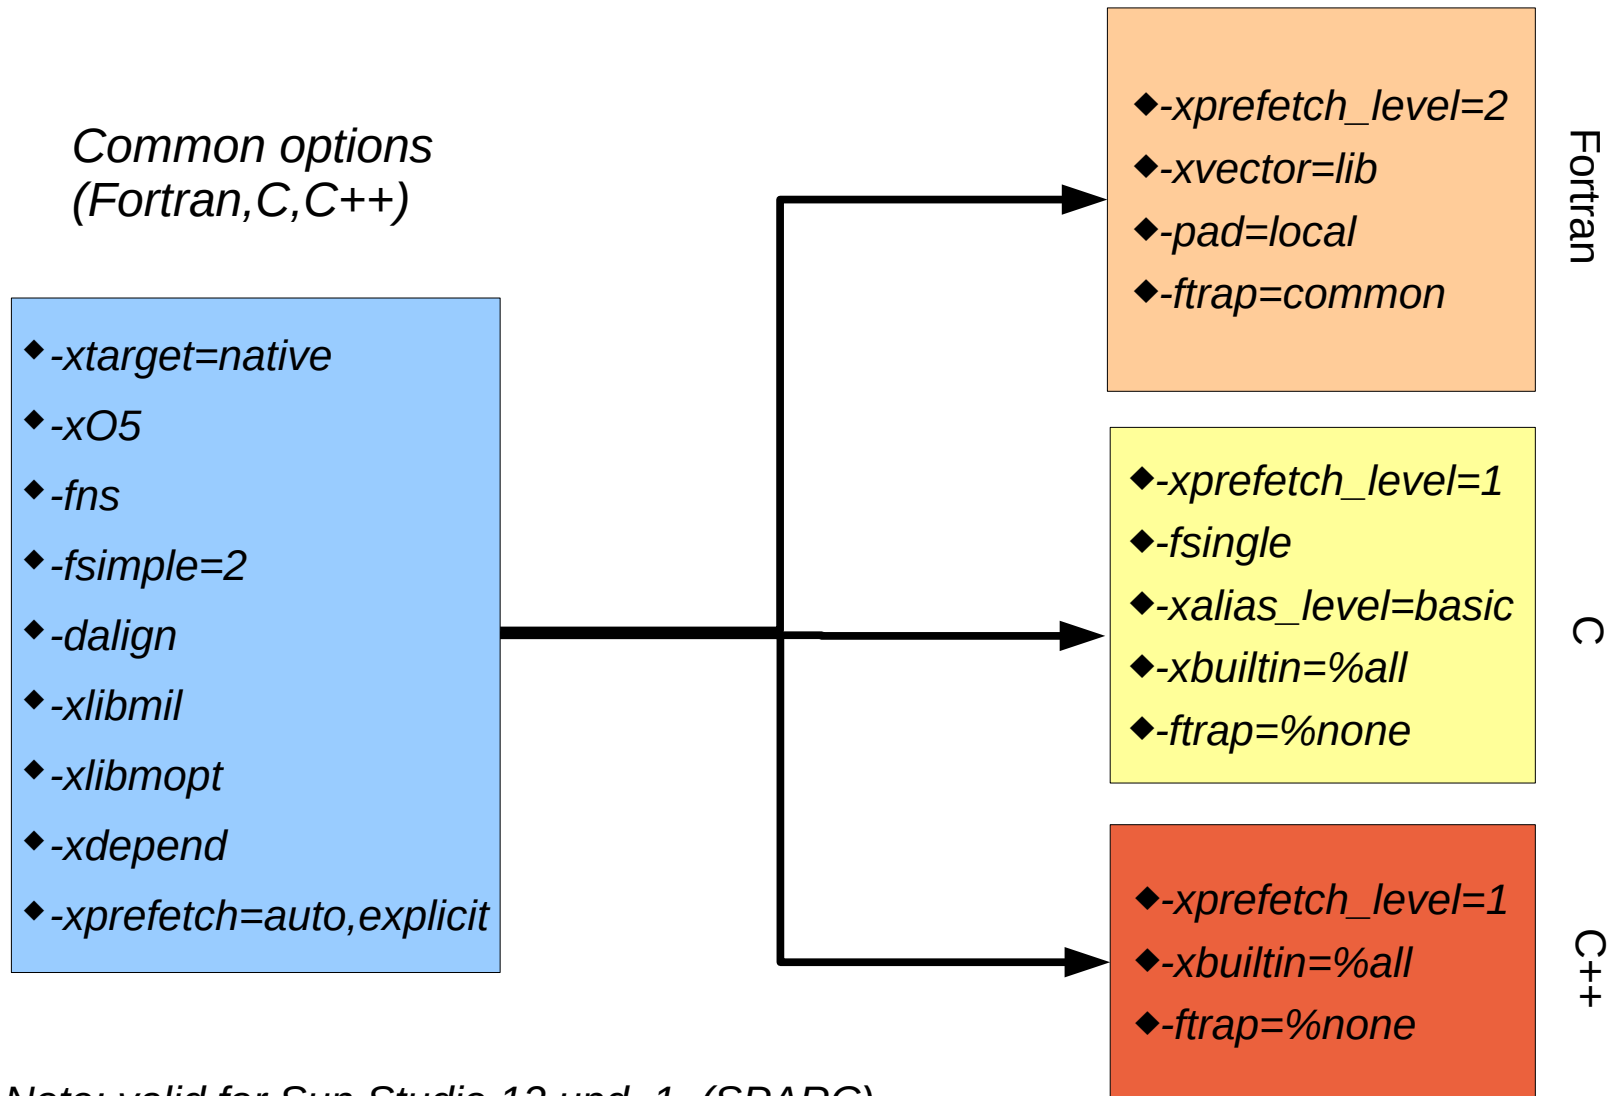
Intel Ivy Bridge: `-g -fast -xchip=ivybridge`

Intel Haswell: `-g -fast -xchip=haswell`

Intel Broadwell: `-g -fast -xchip=broadwell`

Intel Skylake: `-g -fast -xchip=skylake`

Oracle Studio: The -fast macro



Note: valid for Sun Studio 12 upd. 1 (SPARC)

What about other compilers?

- ❑ Most compilers have similar options, that combine many optimizations into a single option
 - ❑ but they do not mean always the same!
- ❑ GCC: -O, -O2, -O3, -Ofast
- ❑ Intel: -O0, -O2 (default!), -O3, -fast
- ❑ look up in the manpages/documentation, what that corresponds to
- ❑ some options have side effects
 - ❑ e.g. Intel and '-fast' changes the linking

GCC: Recommendations

- ❑ a good start: `-g -O3`
- ❑ show optimizer options (“expand -O3”):
 - ❑ `gcc -Q -help=optimizers -O3`
 - ❑ shows a list of all known '-f...' options and their status
- ❑ switch extra options on/off, e.g. loop unrolling
 - ❑ `-funroll-loops` or `-fno-unroll-loops`
- ❑ finding the differences: dump output of command above into files, and run `diff` on the files

GCC: Recommendations

- ❑ some differences between -O3 and -Ofast
 - ❑ e.g. math related options

Option	-O3	-Ofast
-fassociative-math	[disabled]	[enabled]
-ffinite-math-only	[disabled]	[enabled]
-fmath-errno	[enabled]	[disabled]
-freciprocal-math	[disabled]	[enabled]
-ftrapping-math	[enabled]	[disabled]
-funsafe-math-optimizations	[disabled]	[enabled]

GCC: Recommendations

- ❑ Use a Makefile and `make` for compiling/linking:

```
OPT      = -g -O3
ISA      = -mavx2
CHIP     = -march=broadwell
CFLAGS   = $(OPT) $(ISA) $(CHIP)
```

- ❑ *Always start with `-O3` !*
- ❑ Then add other options, to change/increase optimization


GCC: Useful options

- ❑ `-Q --help=optimizers -On` : displays the active optimizer settings for optimization level 'On' (long list)
- ❑ `-Q --help=params` : displays the internal parameters
- ❑ `-fopt-info` : show optimization info at compile time
- ❑ `-v` : displays the configured features
- ❑ `--version` : Shows the compiler version

```
$ gcc --version
gcc (GCC) 6.3.0
Copyright (C) 2016 Free Software Foundation, Inc.
```

Inter Procedural Optimization

- ❑ When used, the compiler stores additional information into the object files
- ❑ This information is used during the link phase to perform additional optimizations



Inter Procedural Optimization

How to use with different compilers?

- ❑ GCC: here it is called 'link time optimization'
 - ❑ -flto
 - ❑ check the documentation for more details
 - ❑ example on next slides
- ❑ clang: uses -flto as well
 - ❑ the internals might differ from GCC
- ❑ Intel:
 - ❑ -ipo

GCC: link time optimization

main.c

```
double init_array(int, float *);

int main(int argc, char *argv[]) {
    int len = atoi(argv[1]);
    float *arr = malloc(len * sizeof(arr));

    // put values into arr
    for(int i = 0; i < len; i++)
        init_array(i, &arr[i]);

    // print the first and last ten values of arr
    for(int i = 0; i < 10; i++)
        printf("a[%d] = %f ... a[%d] = %f\n",
            i, arr[i], len-10+i, arr[len-10+i]);

    return(0);
}
```

GCC: link time optimization

init.c

```
#include <math.h>

double
init_array(int n, float *val) {

    *val = (float)n;

    return sin(n);
}
```

We do not use the
return value in
main, i.e. all calls to
sin() are wasted!

```
$ make
gcc -g -O3 -funroll-loops -c -o main.o main.c
gcc -g -O3 -funroll-loops -c -o init.o init.c
gcc -g -O3 -funroll-loops -o ipo_ex.gcc main.o init.o -lm

$ $ time ./ipo_ex.gcc 100000000 > /dev/null
real    0m4.618s
user    0m4.530s
sys      0m0.088s
```

GCC: link time optimization

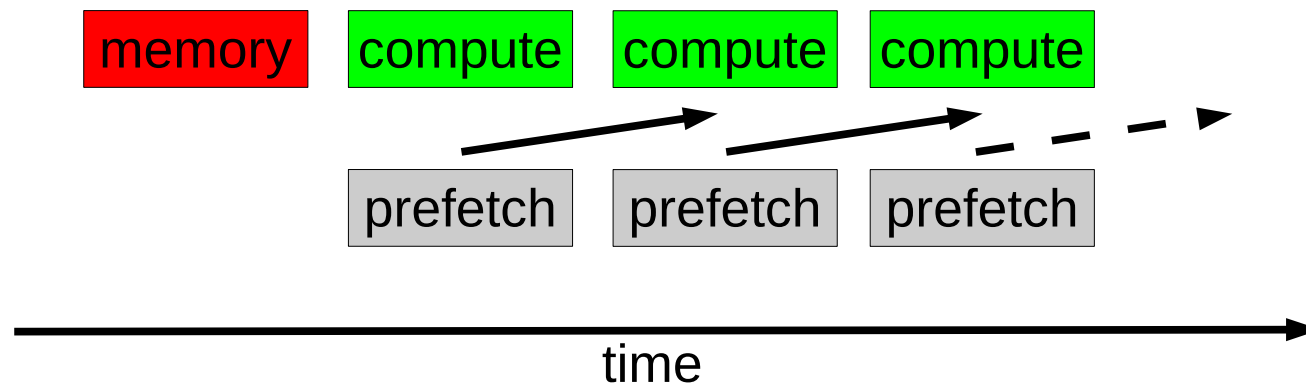
```
$ make
gcc -g -O3 -funroll-loops -flto -c -o main.o main.c
gcc -g -O3 -funroll-loops -flto -c -o init.o init.c
gcc -g -O3 -funroll-loops -flto -o ipo_ex.gcc main.o init.o
-lm

$ time ./ipo_ex.gcc 100000000 > /dev/null
real    0m0.115s
user    0m0.020s
sys     0m0.093s
```

- ❑ in the compile phase, extra information is generated
- ❑ the linker can use this information to optimize further
- ❑ here, the calls to `sin()` are removed, as we do not use the result
- ❑ no reference to `sin()` in the executable

Prefetch: Hiding memory latency

- ❑ The number of clock cycles to access memory increases with the CPU clock speed.
- ❑ Prefetch is a way to overcome this:
 - Fetch data ahead in time, anticipating future use.
- ❑ Special prefetch instructions must be available




Prefetch Support

- ❑ Prefetch is a common feature in modern CPUs: both data and instruction prefetch.
- ❑ Implementation is system dependent!
- ❑ There is both
 - ❑ software prefetch (compiled into the program)
 - ❑ hardware prefetch – which often cannot be disabled
- ❑ x86_64 CPUs have HW prefetch
- ❑ next slide: effect of prefetch on a CPU w/o hardware prefetch

Prefetch: example

Example: Matrix times vector in C (row version)



cache appears to be larger

US-IIIi @ 1062 MHz
L1 : 64 kB
L2 : 1 MB
Peak : 2.1 Gflop/s

GCC: Prefetch options

- ❑ enabled with -O2 and higher
 - ❑ -fprefetch-loop-arrays
 - ❑ there is a number of parameters, that can be tuned
 - ❑ simultaneous-prefetches
 - ❑ min-insn-to-prefetch-ratio
 - ❑ prefetch-min-insn-to-mem-ratio
 - ❑ check with 'gcc -Q --help=params'
 - ❑ usually it is not necessary to change the defaults
 - ❑ has hardly any effect, due to HW prefetch

Pointer overlap – or “aliasing”

```
void vecadd(int n, double *a, double *b, double *c)
{
    for(int i = 0; i < n; i++)
        c[i] = a[i] + b[i];
}
```

```
vecadd(n, &a[0], &b[0], &a[1]);
```

```
void vecadd(n, &a[0], &b[0], &a[1])
{
    for(int i = 0; i < n; i++)
        a[i+1] = a[i] + b[i];
}
```

**Data
dependency!**

Pointer overlap – or “aliasing”

- ❑ Pointer aliasing problem: The C compiler has to assume that different pointers may overlap:
 - ❑ Correct – but non-optimal – code will be generated
 - ❑ Only the programmer might know, that there is no overlap.
- ❑ You can tell the compiler that there is no overlap, using the `restrict` keyword
- ❑ Note: It is then your responsibility that this assumption will not be violated!

No pointer overlap – use of 'restrict'


- ❑ If you can assure that the pointers don't overlap, you can fix your code using the C99 'restrict' keyword:

```
void
vecadd(int n, double * restrict a,
       double * restrict b, double * restrict c)
{
    for(int i = 0; i < n; i++)
        c[i] = a[i] + b[i];
}
```

- ❑ Needs a C99 compliant compiler to be portable!
- ❑ Wrong results, if called with overlapping pointers!

Compiler optimization vs 'restrict'

Vector addition example



compiler
generated
multiple
versions

Faster – but
wrong results!

XeonE5-2650v4 @ 2.2GHz
 L1 : 32 kB
 L2 : 256 kB
 L3 : 30 MB
 GCC 10.2
 -O3 -funroll-loops

Compiler optimization vs 'restrict'

- ❑ compilers can create multiple versions of loops
- ❑ runtime decision, which path to take
 - ❑ no overlap in data: use optimized version
 - ❑ overlap in data: use the “slow” – but correct – version
- ❑ with `'restrict'` keyword:
 - ❑ only the optimized version is generated
 - ❑ gives wrong results, if not called correctly
- ❑ for simple codes like here, leave it to the compiler!

Other tricks

Reconstruct the compiler options from the object files and/or executable:

- ❑ `dwarfdump file.o`
- ❑ and look for keywords
 - ❑ `command_line` or `producer`
- ❑ Very useful to check what has been done to compile the code.

```
$ gcc -g -O3 -funroll-loops -c -o init.o init.c
$ dwarfdump init.o | grep producer
  DW_AT_producer      GNU C17 10.2.0 -mtune=generic
                        -march=x86-64 -g -O3 -funroll-loops
```

Analysis tools

Analysis tools

- ❑ analysis tools are useful to detect bottlenecks in codes
- ❑ modern analysis tools (unlike “old” profilers) work even on 'non-instrumented' code: no need to recompile (in principle)
- ❑ runtime profiles down to the source level (profilers usually work on function/subroutine level)

Analysis tools

- ❑ gprofng
 - ❑ Linux – open source
 - ❑ part of GNU ‘binutils’ since version 2.39
 - ❑ formerly “Oracle Studio Performance Analyzer”
- ❑ ‘perf’ command line tool (Linux)
- ❑ Intel: Vtune Performance Analyzer
 - ❑ Windows/Linux
 - ❑ part of the OneAPI tools (free)
- ❑ Mac OS X: Instruments (part of Xcode)

gprofng – application profiling tool

- ❑ gprofng provides a powerful toolset for runtime application performance analysis
- ❑ both command line tools – and a GUI
- ❑ gprofng works with sub-commands
 - ❑ ... collect app – Command to collect performance data
 - ❑ ... display text – Command to analyze performance data in ASCII format (good for scripting)
 - ❑ ... display gui – GUI for collecting and analyzing performance data

gprofng – application profiling tool

- ❑ Useful on-line documentation for gprofng
 - ❑ Blog:
[gprofng: The Next Generation GNU Profiling Tool](#)
 - ❑ Wiki page: [The gprofng Application Profiling Tool](#)
 - ❑ Video:
<https://www.youtube.com/watch?v=JvnWlv2THTg>
 - ❑ [gprofng GUI](#) - link to the GUI homepage (the GUI is not part of the standard binutils, and might need to be installed as an add-on)

gprofng – application profiling tool

The screenshot displays the gprofng GUI window titled "test.1.er - gprofng GUI (on n-62-30-6)". The interface includes a menu bar (File, Views, Metrics, Tools, Help), a toolbar, and a sidebar with various views (Welcome, Overview, Functions, Timeline, Call Tree, Flame Graph, Source/Disass..., Disassembly, Callers-Callees, Experiments, Threads, Processes, More...). The main area shows a table of CPU time data, and the right sidebar provides selection details for the selected function.

Total CPU Time Table:

EXCLUSIVE				INCLUSIVE				Name
sec.	%	sec.	%	sec.	%	sec.	%	
1.101	100.00	1.101	100.00					<Total>
0.680	61.82	0.680	61.82					distance
0.250	22.73	0.250	22.73					distcheck
0.170	15.45	0.170	15.45					getrusage
0.	0.	0.931	84.55					__libc_start_main
0.	0.	0.931	84.55					main
0.	0.	0.170	15.45					xtime

Selection Details:

- Name: distance
- PC Address: 2:0x000014A0
- Size: 171
- Source File: distance.c
- Object File: st.1.er/archives/aos.gcc_XWsjKAP9gQf}
- Load Object: aos.gcc (found as test.1.er/archives/
- Mangled Name:
- Aliases:

Total CPU Time Summary:

Exclusive		Inclusive	
sec.	%	sec.	%
0.680	61.82%	0.680	61.82%

Called-by / Calls:

distance		distance		
Total CPU Time	distance	Total CPU Time	distance	
ATTRIBUTED	is called by	ATTRIBUTED	calls	
sec.	%	sec.	%	
0.680	100.00			main

Filters: To add a filter, select a row from a view (such as Functions)

Compare: 12.3, 17, 11x, 11x

Status Bar: Local Host: n-62-30-6 | Remote Host: | Working Directory: .../tune_labs_new_sol | Compare: off | Filters: off | Warnings | 2/7

gprofng – application profiling tool

The screenshot displays the gprofng GUI window titled "test.1.er - gprofng GUI (on n-62-30-6)". The interface includes a menu bar (File, Views, Metrics, Tools, Help), a toolbar, and a sidebar with navigation options: Welcome, Overview, Functions, Timeline, Call Tree, Flame Graph, Source/Disass..., Disassembly, Callers-Callees, Experiments, Threads, Processes, and More... The "Source/Disass..." view is active, showing the source code of "distance.c" with a table of CPU times. The function "distance" is highlighted, showing its source code and assembly output.

Total CPU Time		distance.c	
sec.	%		
0.	0.	1.	#include <math.h>
0.	0.	2.	#include "data.h"
0.	0.	3.	
0.	0.	4.	#define sqr(x) ((x)*(x))
0.	0.	5.	
0.	0.	6.	#ifdef ARRAY_OF_STRUCTS
0.	0.	7.	double
0.	0.	8.	<Function: distance>
0.	0.	9.	distance(particle_t *data, int n) {
0.	0.	10.	double dist = 0.0;
0.	0.	11.	double dist2;
0.510	46.36	12.	
0.	0.	13.	for(int i = 0; i < n; i++) {
0.050	4.55	14.	dist2 = sqr(data[i].x);
0.	0.	15.	dist2 += sqr(data[i].y);
0.070	6.36	16.	dist2 += sqr(data[i].z);
0.050	4.55	17.	data[i].dist = sqrt(dist2);
0.	0.	18.	dist += data[i].dist;
0.	0.	19.	}
0.	0.	20.	
0.	0.	21.	return dist;
0.	0.	22.	}

The assembly output for the "distance" function is shown below the source code:

```

[13] 4014a0: test    %esi,%esi
[13] 4014a2: jle     .+0x66 [ 0x401508 ]
[ 8] 4014a4: push   %rbx
[ 8] 4014a5: lea     -0x1(%rsi),%eax

```

The "Selection Details" panel on the right shows the following information:

- Name: distance
- PC Address: 2:0x000014A0
- Size: 171
- Source File: distance.c
- Object File: test.1.er/archives/aos.gcc_XWsjKAP9gQf
- Load Object: aos.gcc (found as test.1.er/archives/)
- Mangled Name:
- Aliases:
- Exclusive: 0.680 { 61.82%}
- Inclusive: 0.680 { 61.82%}

The status bar at the bottom indicates: Local Host: n-62-30-6 | Remote Host: | Working Directory: .../tune_labs_new_sol | Compare: off | Filters: off | Warnings

Hardware Performance Counters

- ❑ Almost all modern CPUs have built-in hardware performance counters:
 - ❑ How many instructions were executed?
 - ❑ How many clock cycles were used?
 - ❑ How many L1 data cache misses occurred?
- ❑ The supported counters are usually listed in the architecture reference manuals.
- ❑ Be aware: The counter names are not for beginners!

Using the Performance Counters

- ❑ Native OS tools, e.g. Linux:

- ❑ perf – Performance monitoring tool

- ❑ requires newer Linux kernel (> 2.6.31)

- ❑ examples:

- `% perf stat -e <event_name> -- command`

- `% perf stat -e <event_name> -p PID`

- `% perf record -e <event_name> -- command`

- `% perf report`

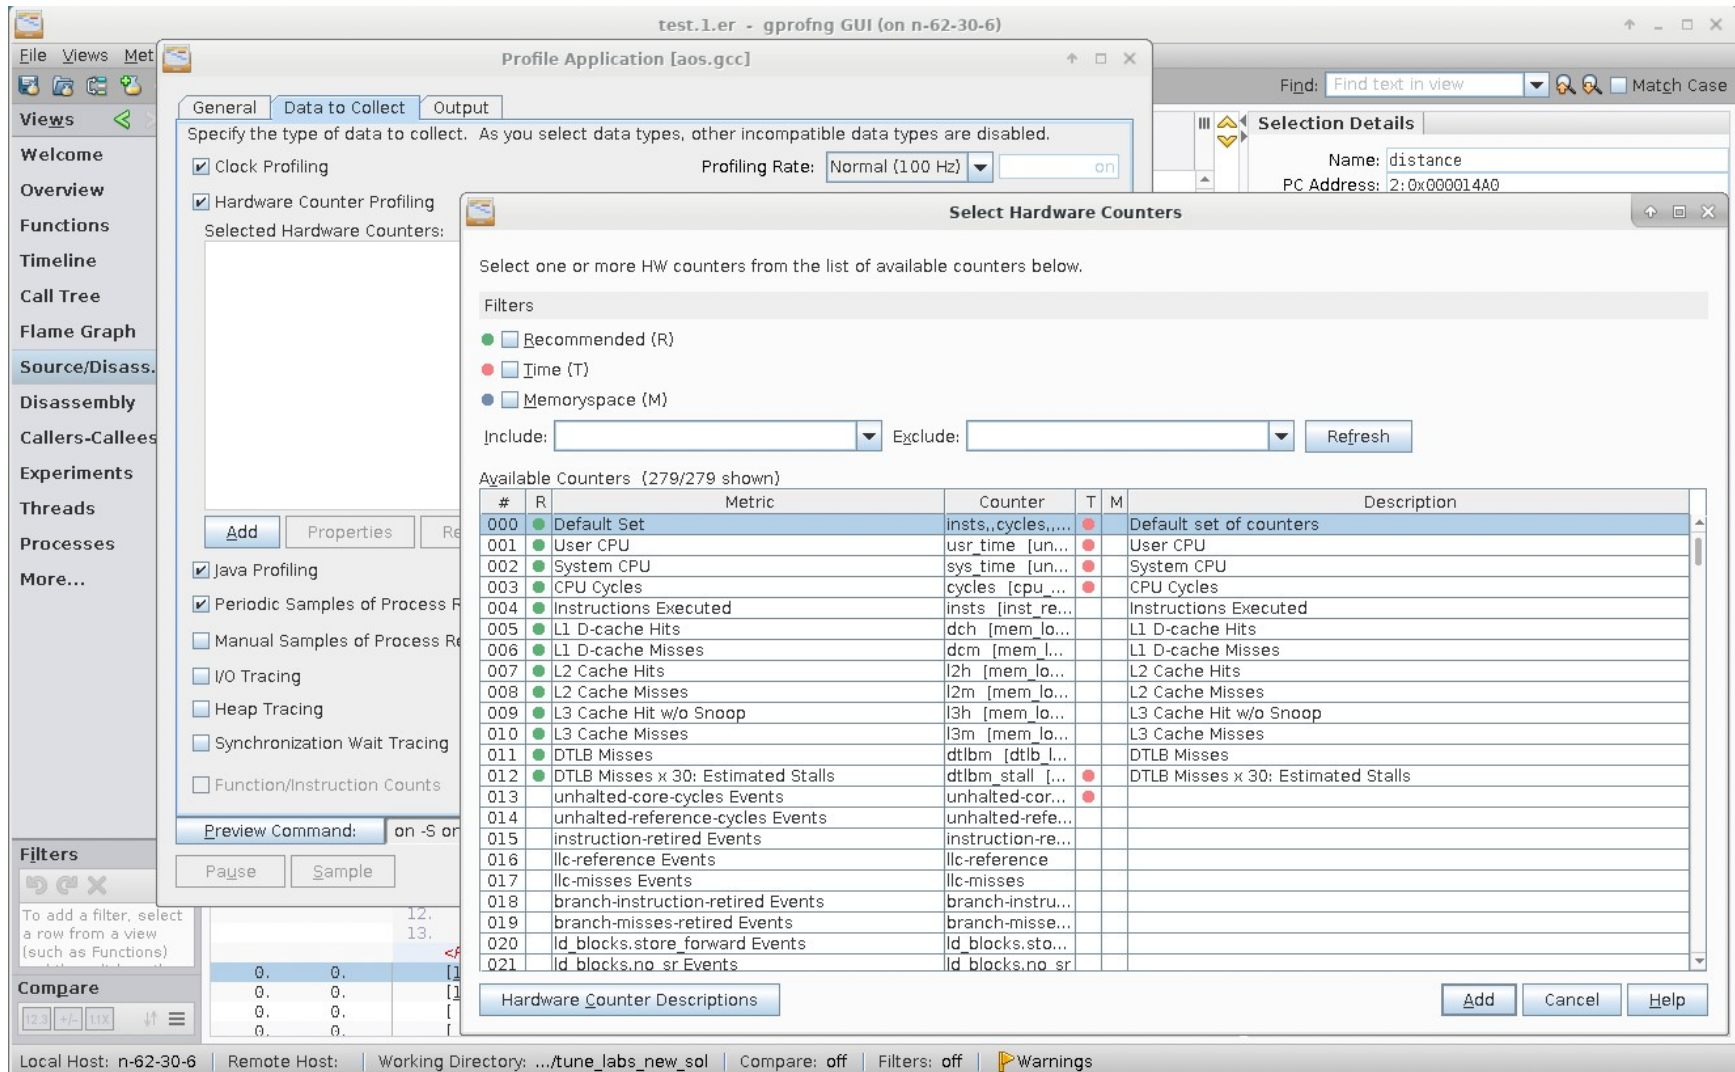
- ❑ 'perf top' - requires root privileges

Using the Performance Counters

- ❑ Available performance counters:
 - ❑ system and CPU dependent
 - ❑ get a list:
 - ❑ `% perf list`
 - ❑ `% gprofng collect app -h`
 - ❑ example: no. of available performance counters on
 - ❑ AMD Opteron: 169
 - ❑ Xeon E5-... v3: 266
 - ❑ Xeon E5-... v4: 311
 - ❑ Xeon Gold... : 246

Using the Performance Counters

Activating performance counters in gprofng GUI:



gprofng – application profiling tool

gprofng demo

Tuning Guide – compact version

- ❑ Make a 'baseline' version (with different data sets/memory requirements)
- ❑ Try to find the best compiler options
 - ❑ with or w/o prefetching
 - ❑ ...
- ❑ Use analysis tools to locate the 'hot spots'
- ❑ Introduce changes: code and/or compiler options
- ❑ Repeat the last two steps until you are satisfied

End of lecture 2